

CS446/ECE 449: Machine Learning

Raymond A. Yeh

yeh17@illinois.edu

University of Illinois at Urbana Champaign, 2020

L10: Deep Learning System & Pytorch

Recap:

- Deep nets & back propagation

Goals of this lecture

- Deep / Machine learning software pipeline
- Pytorch Overview
- Pytorch Example

Reading material

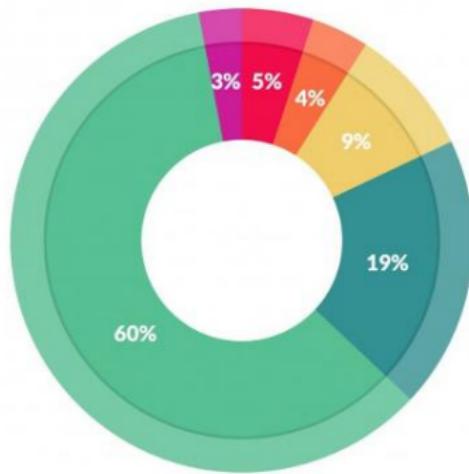
- Pytorch Tutorial (<https://pytorch.org/tutorials/>)

Machine Learning / Deep Learning Systems:

- Data \mathcal{D}
- Feature representation x
- Output representation y
- Model $F(\mathbf{w}, x, y)$
- Loss function L
- Optimization Algorithm

Data & Representation

- How to represent the data?
- Data processing (cleaning data, a large portion of the work.).
- Feature extraction refers to the action of converting from data to $\phi(x)$
- Cleaning data is hard!



What data scientists spend the most time doing

- Building training sets: 3%
- Cleaning and organizing data: 60%
- Collecting data sets: 19%
- Mining data for patterns: 9%
- Refining algorithms: 4%
- Other: 5%

Toy Example Let's say we want to predict whether a student will get an A in ECE544NA. We have data from previous semester,

id	Hours studied	HW grade	Favorite animal	Final grade
1	10	90	cat	99%
2	20	100	elephant	100%
3	0	50	zebra	85%
:	:	:	:	:

- ① How to represent output y ? (Regression $\in \mathbb{R}$? Classification $\in \{-1, 1\}$?)
- ② How to represent input x ? (Recall, $x \in \mathbb{R}^d$.)
Do we need to use all the data?
- ③ Linear regression or logistic regression?

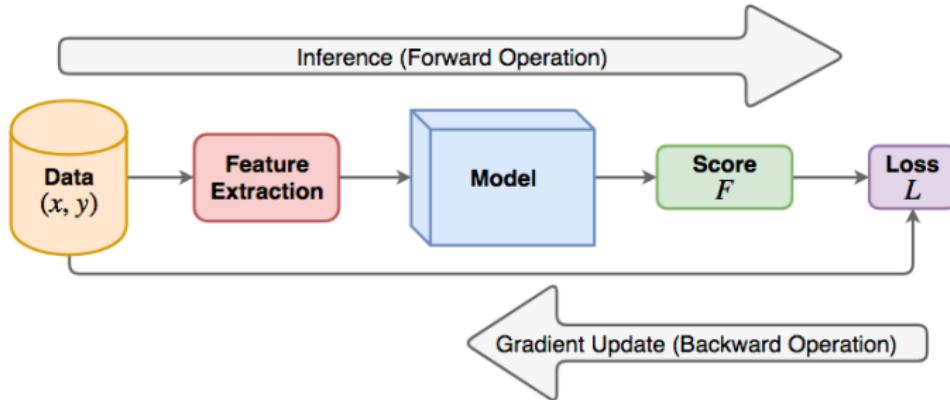
Model & Loss

- Model initialization: initialize all parameters in the model
- Forward operation: generates the score F given an input x
- Prediction operation: takes a score and maps from F to \hat{y}
- Loss: takes F and y and outputs a \mathbb{R}
- Backward operation (Optimize): updates the model parameters to minimize loss

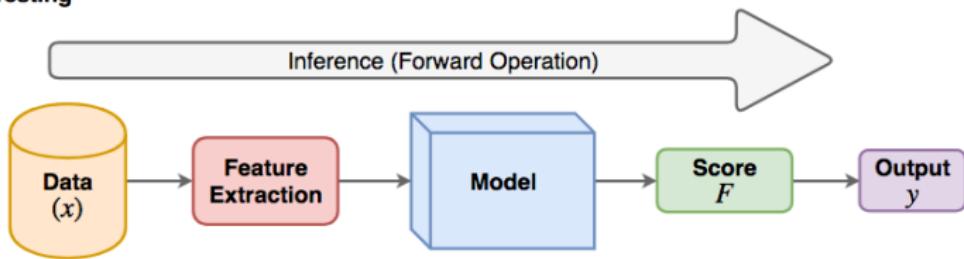
"All models are wrong, but some are useful" — George Box

Train and Test Stage

Training



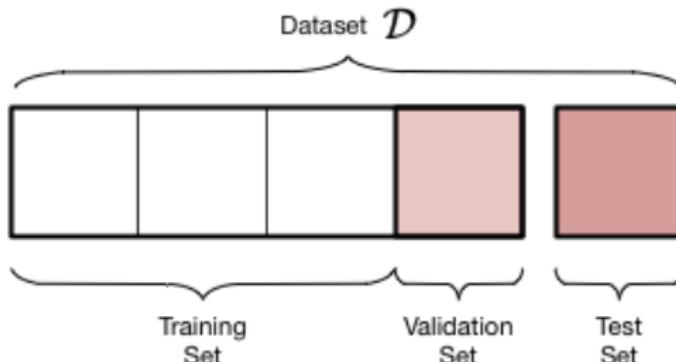
Testing



Machine learning project structure

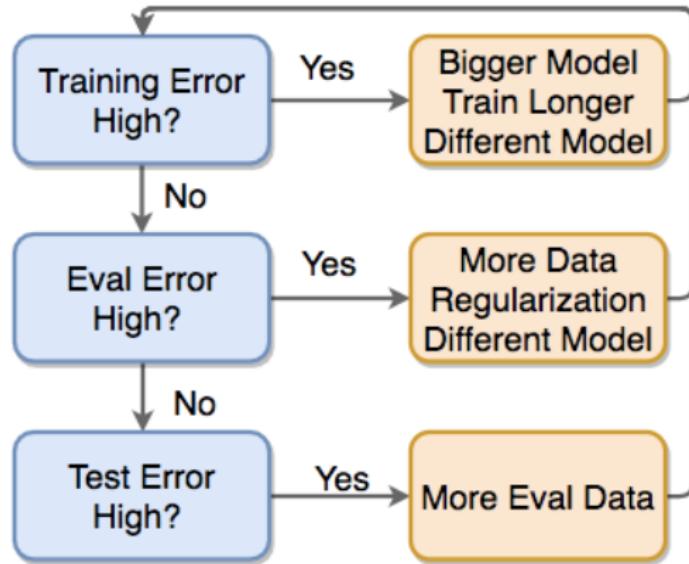
```
data
    processed      <- The data sets for modeling.
    raw           <- Raw data
checkpoints  <- Save and logs experiments
models        <- Model implementations
dataset       <- Dataset implementations
experiment    <- Train and test scripts
main.py
```

Three phases of experiments.



- Training: Update model parameters based on training data
- Evaluation: Measure performance on the validation data to choose the hyperparameters
- Testing: Measure performance on the testing data

A standard machine learning workflow.



Credit: Andrew Ng

Why is machine learning hard?

- Data is noisy
- Many possibilities for data representation
- Hard to debug (many things could go wrong!)



Pytorch

-  PyTorch
- Facebook's deep learning framework.

TensorFlow: machine learning library for research and production



-  TensorFlow
- Google's deep learning framework.

Others

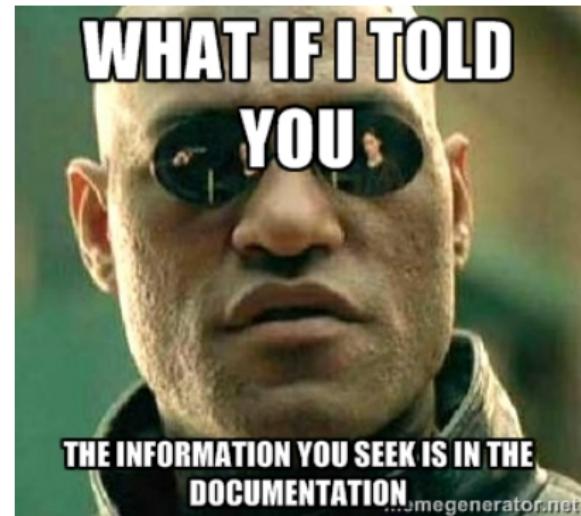
- MXNET: A flexible and efficient library for deep learning
- JAX: Autograd and XLA

Pytorch

- Installation
(<https://pytorch.org/get-started/locally/>)
- Using Google Colab (CPU and GPUs for Free!).

A few useful Pytorch Packages

- torch
- torch.nn
- torch.optim
- torchvision



Data with Pytorch

- Tensor \approx multidimensional array.
- First dimension is, conventionally, used for batching.

```
1 import torch
2 num_example = 3
3 feature_dim = 2
4 x = torch.randn(num_example, feature_dim)
5 print(x)
```

```
tensor([[-1.6828,  1.4117],
       [ 0.0263, -0.8230],
       [ 1.2949, -1.2613]])
```

Mathematical operations with Pytorch

- Matrix multiplication (torch.mm, torch.matmul)

```
1 | import torch
2 | x = torch.randn(3, 3)
3 | out = torch.mm(x, x.t())
```

- Pseudo-inverse (torch.pinverse)
- Least squares solver (torch.lstsq)
- ...

Model with Pytorch

- Represent a model via a computation graph (acyclic)
- Standard building block ("modules") for deep-nets are in `torch.nn`
- Standard loss functions are also in `torch.nn`

```
1  class LinearClassifier(nn.Module):  
2      # Adapted from torch.nn.Linear  
3      def __init__(self, in_features=2, out_features=1):  
4          super(LinearClassifier, self).__init__()  
5          self.w = nn.Parameter(torch.Tensor(out_features,  
6                                              in_features))  
7          self.b = nn.Parameter(torch.Tensor(out_features))  
8  
9      def forward(self, input):  
10         x = F.linear(input, self.w, self.b)  
11         return torch.sigmoid(x)  
12  
13     def predict(self, input):  
14         return F.linear(input, self.w, self.b) > 0  
15     ...
```

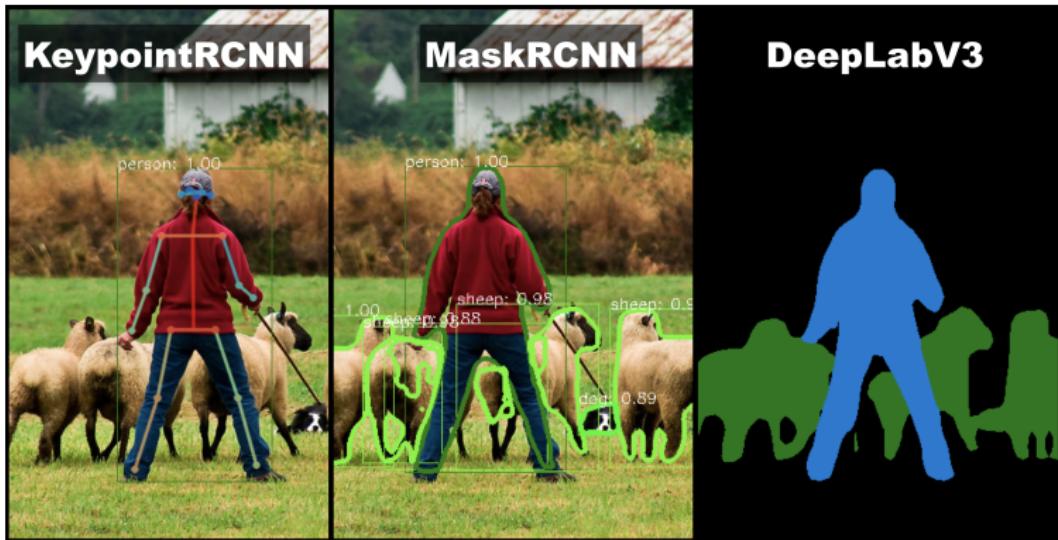
Optimizer

- Standard gradient-based optimizers are in `torch.optim`

```
1 # Training
2 model = LinearClassifier(in_features=feature_dim)
3 epochs = 100
4 optimizer = torch.optim.SGD(model.parameters())
5 loss_epoch = []
6 acc_epoch = []
7 for epoch in range(epochs):
8     score = model(x)
9     optimizer.zero_grad()
10    loss = model.loss(score, y.float().unsqueeze(-1))
11    loss.backward()
12    optimizer.step()
```

Computer Vision (torchvision)

- `torchvision.datasets`
- `torchvision.transforms`
- `torchvision.models`



Credit: Pytorch

Demo on Google Colab [[Link](#)]