

# CS 446/ECE 449: Machine Learning

A. G. Schwing

University of Illinois at Urbana-Champaign, 2020

## L26: Generative Adversarial Nets

## Goals of this lecture

- Getting to know Generative Adversarial Nets
- Understanding generative methods
- Differentiating between discriminative and generative methods

## Reading Material:

- I. Goodfellow et al.; Generative Adversarial Networks; [arxiv.org/abs/1406.2661](https://arxiv.org/abs/1406.2661)
- M. Arjovsky et al.; Wasserstein GAN; [arxiv.org/abs/1701.07875](https://arxiv.org/abs/1701.07875)

**Recap:** Maximum likelihood so far?

Model:

$$p(\mathbf{y}|\mathbf{x}) = \frac{\exp F(\mathbf{y}, \mathbf{x}, \mathbf{w})/\epsilon}{\sum_{\hat{\mathbf{y}}} \exp F(\hat{\mathbf{y}}, \mathbf{x}, \mathbf{w})/\epsilon}$$

- $\mathbf{y}$ : discrete output space
- $\mathbf{x}$ : input data

**Now:**

How about modeling a distribution  $p(\mathbf{x})$  for the data?

Why modeling a distribution  $p(x)$ ?

- Synthesis of objects (images, text)
- Environment simulator (reinforcement learning, planning)
- Leveraging unlabeled data


Given data points  $x$ , how can we model  $p(x)$ ?

- Maximum likelihood approach:

$$\theta^* = \max_{\theta} \sum_i \log p(x^{(i)}; \theta)$$

- ▶ Fit mean and variance (= parameters  $\theta$ ) of a distribution (e.g., Gaussian)
- ▶ Fit parameters  $\theta$  of a mixture distribution (e.g., mixture of Gaussian, k-means)
- ▶ Use a variational auto-encoder

## Variational auto-encoder:



## Loss function:

$$\mathcal{L}(p_{\theta}, q_{\phi}) \approx -D_{KL}(q_{\phi}, p) + \frac{1}{N} \sum_{i=1}^N \log p_{\theta}(x|z^i)$$

## Issue that we had to address:

Computing a normalization constant

Another approach:

## Generative Adversarial Nets



Main idea:

Don't write a formula for  $p(x)$ , just learn to sample directly

Advantage:

No summation over large probability spaces

Formulate the problem as a game between two players:

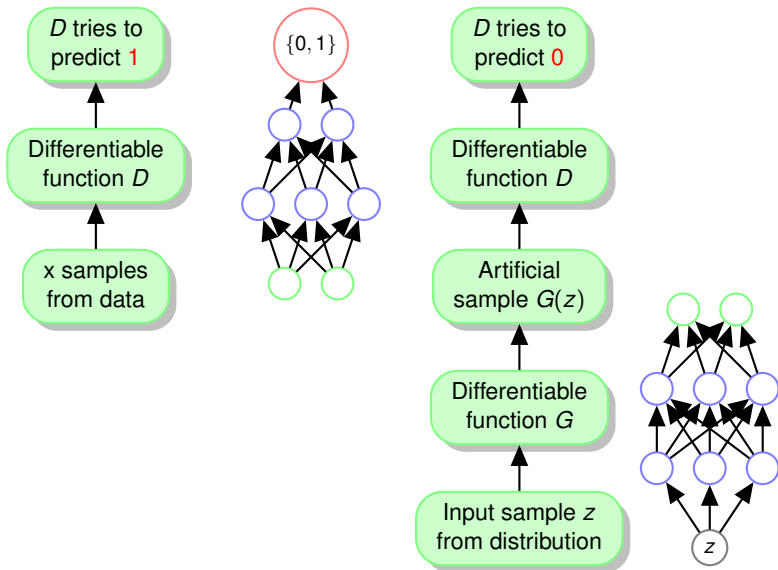
- Generator  $G$
- Discriminator  $D$

Task of the players

- $G$  generates examples
- $D$  predicts whether the example is artificial or real

$G$  tries to “trick”  $D$  by generating samples that are hard to distinguish

## Pictorially:



Mathematically:

- Generator  $G_\theta(z)$
- Discriminator  $D_w(x) = p(y = 1|x)$

How to choose  $w$ :

$$\min_w - \sum_x \log D_w(x) - \sum_z \log(1 - D_w(G_\theta(z)))$$

How to choose  $\theta$ :

$$\max_\theta \min_w - \sum_x \log D_w(x) - \sum_z \log(1 - D_w(G_\theta(z)))$$

Generative adversarial nets:

$$\max_{\theta} \min_w - \sum_x \log D_w(x) - \sum_z \log(1 - D_w(G_{\theta}(z)))$$

How to optimize this theoretically?

Repeat until stopping criteria

- 1 Gradient step w.r.t.  $w$
- 2 Gradient step w.r.t.  $\theta$

In practice:

Heuristics make this optimization more stable.

Analysis of generative adversarial nets:

What is the optimal discriminator assuming arbitrary capacity?

$$\begin{aligned}\min_D : \quad & - \int_x p_{\text{data}}(x) \log D(x) dx - \int_z p_z(z) \log(1 - D(G_\theta(z))) dz \\ & = - \int_x p_{\text{data}}(x) \log D(x) + p_G(x) \log(1 - D(x)) dx\end{aligned}$$

Euler-Lagrange formalism:

$$S(D) = \int_x L(x, D, \dot{D}) dx$$

Stationary  $D$  from

$$\frac{\partial L(x, D, \dot{D})}{\partial D} - \frac{\cancel{d}}{\cancel{dx}} \frac{\partial L(x, D, \dot{D})}{\partial \dot{D}} = 0$$

What is the optimal discriminator assuming arbitrary capacity?

$$\frac{\partial L(x, D, \dot{D})}{\partial D} = -\frac{p_{\text{data}}}{D} + \frac{p_G}{1-D} = 0$$

Consequently:

$$D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)}$$

Given the optimal discriminator, what is the optimal generator?

$$\begin{aligned} & - \int_x p_{\text{data}}(x) \log D^*(x) + p_G(x) \log(1 - D^*(x)) dx \\ = & - \int_x p_{\text{data}}(x) \log \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)} + p_G(x) \log \frac{p_G(x)}{p_{\text{data}}(x) + p_G(x)} dx \\ = & - 2 \text{JSD}(p_{\text{data}}, p_G) + \log(4) \end{aligned}$$

$$\text{JSD}(p_{\text{data}}, p_G) = \frac{1}{2} \text{KL}(p_{\text{data}}, M) + \frac{1}{2} \text{KL}(p_G, M) \quad \text{with} \quad M = \frac{1}{2}(p_{\text{data}} + p_G)$$

Consequently:

$$p_G(x) = p_{\text{data}}(x)$$





Some heuristics that improve optimization of:

$$\max_{\theta} \min_w - \sum_x \log D_w(x) - \sum_z \log(1 - D_w(G_{\theta}(z)))$$

- If  $G$  is very bad and  $D$  is very good: almost no gradient
- Solve instead


$$\min_{\theta} - \sum_z \log D_w(G_{\theta}(z))$$

- Issue: no joint cost function for  $D$  and  $G$




Plenty of additional impressive tricks.


## Some results on toy data:




0 Iterations



2000 Iterations




4000 Iterations



5000 Iterations

## Application: modeling of images



















Text description	This flower has petals that are white and has pink shading	This flower has a lot of small purple petals in a dome-like configuration	This flower has long thin yellow petals and a lot of yellow anthers in the center	This flower is pink, white, and yellow in color, and has petals that are striped	This flower is white and yellow in color, with petals that are wavy and smooth	This flower has upturned petals which are thin and orange with rounded edges	This flower has petals that are dark pink with white edges and pink stamen
64x64 GAN-INT-CLS [22]							
256x256 StackGAN							

Figure 4. Example results by our proposed StackGAN and GAN-INT-CLS [22] conditioned on text descriptions from Oxford-102 test set.

## Quiz:

- What generative modeling techniques do you know about?
- What are the short-comings of those techniques?
- What differentiates GANs from other generative models?
- What are the short-comings of GANs.

## Important topics of this lecture

- Getting to know generative adversarial nets (GANs)
- Understanding their advantages and disadvantages

## Next up:

Other image generation/completion techniques